

Phan Thanh Tan

Data Engineer Intern

✉ thanhtanphan011022@gmail.com | ☎ +84 984 651 311 | 🌐 Tan Thanh

📍 Phan Thanh Tan | Ho Chi Minh, Viet Nam

SUMMARY

Proactive and data-driven fourth year Data Engineering student with a background in ETL, SQL, Python and hands-on experience building data pipelines using open source tools. Seeking a Data Engineer Intern/Junior position to contribute to scalable data solutions while fast-tracking my technical growth in a professional production environment.

EDUCATION

Ho Chi Minh University of Technology

Sep. 2023 – Jun. 2027 (Expected)

Bachelor of Computer Science, Major: Data Engineering and Big Data

Ho Chi Minh City, Vietnam

- **Relevant Coursework:** Data Engineering, Machine Learning, Artificial Intelligence, Database Management System, Data Mining, Big Data Analysis, Data Warehouse, Cloud Computing, Object Oriented Programming (OOP).

TECHNICAL STACKS

Programming Languages: Python, C++, R, SQL, TypeScript, Javascript

Tools & Cloud: PySpark, MLflow, FastAPI, Docker, MinIO/S3, Apache Airflow, dbt, Google BigQuery, Kafka

Libraries & Others: Streamlit, Power BI, scikit-learn, Pandera, Alembic, GitHub Actions

EXPERIENCE

MyTegami.win

May 2025 – Jul. 2025

Backend Developer Intern

Remote

Designed and engineered the backend AI generation flow for a bilingual cover letter service using Genkit and Gemini 2.0 Flash, defining strict Zod schemas for structured outputs. Bring the 14 languages, 5 tone options, and revenue generation to the forefront as the concrete result of your engineering.

PROJECTS

Scalable Customer Data Platform (CDP) | *Personal Project*

[View Project](#)

Tech Stack: PySpark, MLflow, FastAPI, Docker Compose, Airflow, PostgreSQL, Python

Oct. 2025 – Feb. 2026

- Designed and built a production-style MLOps platform covering the full ML lifecycle: PySpark ETL → MLflow tracking → FastAPI serving → A/B testing → Streamlit BI dashboard.
- Implemented CI/CD with GitHub Actions, database schema migrations with Alembic, and a hybrid execution mode bridging S3/MinIO and local Parquet storage.
- Achieved **93.5% churn prediction accuracy** (ROC-AUC 0.99) on a 7,043-record dataset, architected to scale to **50M users** via distributed Spark processing.

Sales Dashboard — Enterprise Data Platform | *Personal Project*

[View Project](#)

Tech Stack: Python, SQL, BigQuery, dbt, Apache Airflow, Docker, Power BI, Streamlit

Dec. 2025 – Mar. 2026

- **Designed a cloud-native ELT pipeline ingesting** Amazon e-commerce data into BigQuery, applying JSON Schema validation and automated data quality gates (zero data inconsistency).
- Modeled a Star Schema using dbt's Medallion Architecture (Bronze → Silver → Gold), reducing query scan costs by **~40–50%** and dashboard loading times by **~40%**.
- **Orchestrated the full lifecycle with Airflow**, enabling 15-minute refresh intervals for near real-time reporting.

Real-time Sentiment Analysis Pipeline | *Personal Project*

[View Project](#)

Tech Stack: Kafka, Spark Structured Streaming, MongoDB, FastAPI, MLflow

Nov. 2025 – Feb. 2026

- Built an end-to-end streaming data pipeline ingesting social media data via Kafka, processing it with Spark Structured Streaming, and serving insights through FastAPI.
- Optimized ML inference throughput by 478x (from ~71 to ~33,300 records/sec) by **replacing Scalar UDFs with Pandas UDFs (Apache Arrow)**.
- Implemented Pydantic schema validation to reject malformed records and integrated MLflow for experiment tracking of a Scikit-Learn model **trained on 30+ emotion categories**.

LANGUAGES & CERTIFICATES

Languages: English (TOEIC 590 - [View Certificate](#)) | Vietnamese (Native)

Certificates: HackerRank Software Engineer Certificate